


RAFT CONSENSUS ALGORITHM

Miulescu Cristina-Maria

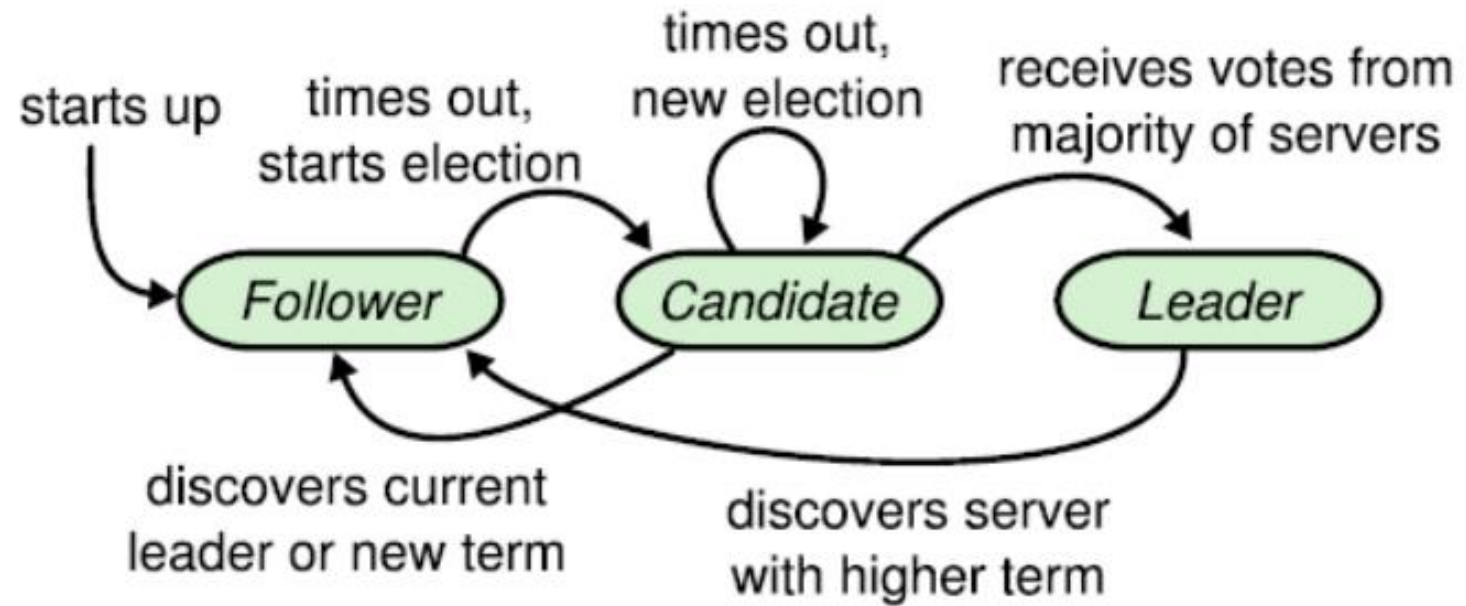
SCPD

WHAT IS RAFT?

- Distributed Consensus Algorithm
- Alternative to Paxos
- Leader-based
- Ensures log replication



MAIN IDEA



IMPLEMENTATION

Server-Client

Log Manager

Raft Implementation



SERVER- CLIENT

```
$ python client.py 10000 1  
[*] Connecting to localhost port 10000
```

```
[*] Type your message:
```

```
>> get ana
```

```
[*] Sending: get ana
```

```
[*] Received:
```

```
>> GET response:
```

```
4
```

```
[*] Type your message:
```

```
>> set d 8
```

```
[*] Sending: set d 8
```

```
[*] Received:
```

```
>> SET response:
```

```
d set to 8
```

```
[*] Type your message:
```

```
>> get d
```

```
[*] Sending: get d
```

```
[*] Received:
```

```
>> GET response:
```

```
8
```

LOG MANAGER

Execute

- Execute commands from the client:
 - Get
 - Set
 - Delete
 - Show

Keep

- Keep track of each node's logs

RAFT IMPLEMENTATION

APPEND ENTRIES RPC

Leader:

- Heartbeat messages
- Replicate log entries

- ❖ Current term
- ❖ Previous term in log
- ❖ Previous index in log

REQUEST VOTE RPC

Candidate:

- Request votes

- ❖ Current term
- ❖ Previous term in log
- ❖ Previous index in log

FOLLOWER

RECOVER LOGS

```
$ python server.py s1 10000  
C:\Users\crist\Documents\Projects\DA\raft\src\logs\server-configuration.txt  
[*] Recovering logs...  
[*] Current term: 6  
[*] Server started on ('localhost', 10000) with timeout 15.0 seconds
```

RECEIVE APPEND ENTRIES RPC

```
GiveEntriesAfterIndex 1  
Sending message  
[*] Connection from ('127.0.0.1', 58804)  
[*] Received from s1 : AppendEntries 6-9-5-[[{'index': '2', 'term': '3', 'c  
ommand': 'get b'}, {'index': '3', 'term': '3', 'command': 'set b 5'}, {'in  
dex': '4', 'term': '3', 'command': 'get b'}, {'index': '5', 'term': '4', '  
command': 'get ana'}, {'index': '6', 'term': '5', 'command': 'get ana'}, {'  
'index': '7', 'term': '5', 'command': 'set d 8'}, {'index': '8', 'term': '  
5', 'command': 'get d'}, {'index': '9', 'term': '5', 'command': 'get d'}]  
Last log: 1 2 set b 30
```


CANDIDATE

```
crist@cristina MINGW64 ~/Documents/Projects/DA/raft/src (raft-dev-2)
$ python server.py s1 10000
C:\Users\crist\Documents\Projects\DA\raft\src\logs\server-configuration.txt
[*] Recovering logs...
[*] Current term: 6
[*] Server started on ('localhost', 10000) with timeout 15.0 seconds
True
[*] Starting election...
Sending message
[*] Connection from ('127.0.0.1', 60971)
[*] Received from s2 : Count on me
[*] I am Leader
{'s1': True, 's2': True, 's3': False, 's4': False}
Sending heartbeat...
Sending message
Sending message
[*] Connection from ('127.0.0.1', 60975)
[*] Received from s3 : Count on me
[*] I am Leader
{'s1': True, 's2': True, 's3': True, 's4': False}
```

LEADER

HANDLE CLIENT CONNECTIONS

```
[*] Received from client : get ana  
Sending heartbeat...  
Sending message  
Sending message  
Sending message  
█
```

MAINTAIN LOG REPLICATION
THROUGH HEARTBEAT

```
[*] Received from s4 : GiveEntriesAfterIndex 1  
[{'index': '2', 'term': '3', 'command': 'get b'}, {'index': '3', 'term': '3', 'command': 'set b 5'}, {'index': '4', 'term': '3', 'command': 'get b'}, {'index': '5', 'term': '4', 'command': 'get ana'}, {'index': '6', 'term': '5', 'command': 'get ana'}, {'index': '7', 'term': '5', 'command': 'set d 8'}, {'index': '8', 'term': '5', 'command': 'get d'}, {'index': '9', 'term': '5', 'command': 'get d'}]
```

WHEN LEADER FAILS...

```
<class 'str'>  
[*] Connection from ('127.0.0.1', 61227)  
[*] Received from s1 : AppendEntries 7-11-6-[]  
Last log: 9 5 get d
```

```
<class 'str'>  
True  
[*] Starting election...  
█
```

WHEN CLIENT CONNECTS TO A FOLLOWER...

```
crist@cristina MINGW64 ~/Documents/Projects/DA/raft/src (raft-dev-2)
$ python client.py 10001 1
[*] Connecting to localhost port 10001

-----
[*] Type your message:
    >> get ana
[*] Sending: get ana
[*] Received:

Sorry, I am not the leader. Last leader I heard from is: s1

-----
[*] Type your message:
    >>
[*] Closing socket
```

MULTIPLE CLIENTS

```
crist@cristina MINGW64 ~/Documents/Projects/DA/raft/src (raft-dev-2)
$ python client.py 10000 1
[*] Connecting to localhost port 10000

-----
[*] Type your message:
    >> get ana
[*] Sending: get ana
[*] Received:

>> GET response:
4

-----
[*] Type your message:
    >> set d 8
[*] Sending: set d 8
[*] Received:

>> SET response:
d set to 8

-----
[*] Type your message:
    >> get d
[*] Sending: get d
[*] Received:

>> GET response:
8

-----
[*] Type your message:
    >> 
```

```
crist@cristina MINGW64 ~/Documents/Projects/DA/raft/src (raft-dev-2)
$ python client.py 10000 2
[*] Connecting to localhost port 10000

-----
[*] Type your message:
    >> get d
[*] Sending: get d
[*] Received:

>> GET response:
8

-----
[*] Type your message:
    >> 
```



CONCLUSION

- Main goal of the Raft Consensus algorithm is understandability which has been easily achieved during the implementation

THANK YOU